

Charlie Crisp

Building a Blockchain Library for OCaml

Computer Science Tripos – Part II

Pembroke College

December 30, 2017

Proforma

Name: **Charlie Crisp**
College: **Pembroke College**
Project Title: **Building a Blockchain Library for OCaml**
Examination: **Computer Science Tripos – Part II, July 2018**
Word Count: **????¹ (well less than the 12000 limit)**
Project Originator: **KC Sivaramakrishnan**
Supervisor: **KC Sivaramakrishnan**

Original Aims of the Project

To build a library in OCaml, which can be used as a building block for Blockchain applications. The library should allow participating nodes to own a shared copy of a Blockchain data structure, agreed upon using consensus. Nodes should also be able to commit transactions to the blockchain, which should then be visible to other participating nodes.

Work Completed

All that has been completed appears in this dissertation.

Special Difficulties

None

¹This word count was computed by `detex diss.tex | tr -cd '0-9A-Za-z \n' | wc -w`

Declaration

I, Charlie Crisp of Pembroke College, being a candidate for Part II of the Computer Science Tripos [or the Diploma in Computer Science], hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed [signature]

Date [date]

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 9 |
| 1.1 | The History of the Blockchain | 9 |
| 1.2 | Building the document | 9 |
| 1.2.1 | The makefile | 10 |
| 1.3 | Counting words | 10 |
| 2 | Preparation | 11 |
| 3 | Implementation | 13 |
| 3.1 | Verbatim text | 13 |
| 3.2 | Tables | 14 |
| 3.3 | Simple diagrams | 14 |
| 3.4 | Adding more complicated graphics | 14 |
| 4 | Evaluation | 17 |
| 4.1 | Printing and binding | 17 |
| 4.2 | Further information | 17 |
| 5 | Conclusion | 19 |
| | Bibliography | 19 |
| A | Latex source | 21 |
| A.1 | diss.tex | 21 |
| A.2 | proposal.tex | 21 |
| B | Makefile | 23 |
| B.1 | makefile | 23 |
| B.2 | refs.bib | 23 |
| C | Project Proposal | 25 |

List of Figures

| | | |
|-----|---|----|
| 3.1 | A picture composed of boxes and vectors. | 14 |
| 3.2 | A diagram composed of circles, lines and boxes. | 15 |
| 3.3 | Example figure using encapsulated postscript | 15 |
| 3.4 | Example figure where a picture can be pasted in | 16 |
| 3.5 | Example diagram drawn using <code>xfig</code> | 16 |
| C.1 | A typical blockchain structure [3] | 2 |

Acknowledgements

This document owes much to an earlier version written by Simon Moore [?]. His help, encouragement and advice was greatly appreciated.

Chapter 1

Introduction

1.1 The History of the Blockchain

The blockchain, in its simplest form, is a series of blocks of data, where each block contains the cryptographic hash of the previous block in the chain. This means that the chain can exhibit arbitrary branching.

The significance of blocks containing hashes of previous blocks, is that if a change were to be made to a previous block, it would change the hashes of all subsequent blocks in the chain. This means that once a block is committed to the blockchain, it cannot be modified.

1.2 Building the document

This document was produced using $\text{\LaTeX} 2_{\epsilon}$ which is based upon $\text{\LaTeX}[?]$. To build the document you first need to generate `diss.aux` which, amongst other things, contains the references used. This is done by executing the command:

```
pdflatex diss
```

Then the bibliography can be generated from `refs.bib` using:

```
bibtex diss
```

Finally, to ensure all the page numbering is correct run `pdflatex` on `diss.tex` until the `.aux` files do not change. This usually takes 2 more runs.

1.2.1 The makefile

To simplify the calls to `pdflatex` and `bibtex`, a makefile has been provided, see Appendix B.1. It provides the following facilities:

```
make
```

Display help information.

```
make proposal.pdf
```

Format the proposal document as a PDF.

`make view-proposal`

Run `make proposal.pdf` and then display it with a Linux PDF viewer (preferably “okular”, if that is not available fall back to “evince”).

`make diss.pdf`

Format the dissertation document as a PDF.

`make count`

Display an estimate of the word count.

`make all`

Construct `proposal.pdf` and `diss.pdf`.

`make pub`

Make `diss.pdf` and place it in my `public.html` directory.

`make clean`

Delete all intermediate files except the source files and the resulting PDFs. All these deleted files can be reconstructed by typing `make all`.

1.3 Counting words

An approximate word count of the body of the dissertation may be obtained using:

```
wc diss.tex
```

Alternatively, try something like:

```
detex diss.tex | tr -cd '0-9A-Z a-z\n' | wc -w
```

Chapter 2

Preparation

This chapter is empty!

Chapter 3

Implementation

3.1 Verbatim text

Verbatim text can be included using `\begin{verbatim}` and `\end{verbatim}`. I normally use a slightly smaller font and often squeeze the lines a little closer together, as in:

```
GET "libhdr"

GLOBAL { count:200; all  }

LET try(ld, row, rd) BE TEST row=all
THEN count := count + 1
ELSE { LET poss = all & ~(ld | row | rd)
      UNTIL poss=0 DO
      { LET p = poss & -poss
        poss := poss - p
        try(ld+p << 1, row+p, rd+p >> 1)
      }
    }
LET start() = VALOF
{ all := 1
  FOR i = 1 TO 12 DO
  { count := 0
    try(0, 0, 0)
    writef("Number of solutions to %i2-queens is %i5*n", i, count)
    all := 2*all + 1
  }
  RESULTIS 0
}
```

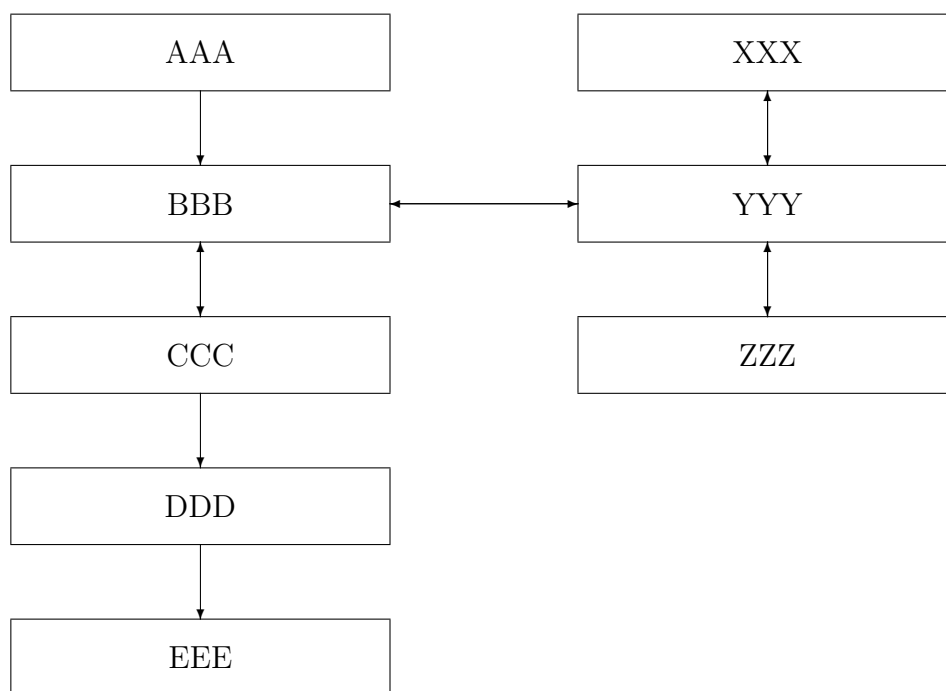


Figure 3.1: A picture composed of boxes and vectors.

3.2 Tables

Here is a simple example¹ of a table.

| Left Justified | Centred | Right Justified |
|-------------------|---------|--------------------|
| First | A | XXX |
| Second | AA | XX |
| Last | AAA | X |

There is another example table in the proforma.

3.3 Simple diagrams

Simple diagrams can be written directly in \LaTeX . For example, see figure 3.1 on page 14 and see figure 3.2 on page 15.

3.4 Adding more complicated graphics

The use of \LaTeX format can be tedious and it is often better to use encapsulated postscript (EPS) or PDF to represent complicated graphics. Figure 3.3 and 3.5 on page 16 are

¹A footnote



Figure 3.2: A diagram composed of circles, lines and boxes.

examples. The second figure was drawn using `xfig` and exported in `.eps` format. This is my recommended way of drawing all diagrams.



Figure 3.3: Example figure using encapsulated postscript

Figure 3.4: Example figure where a picture can be pasted in

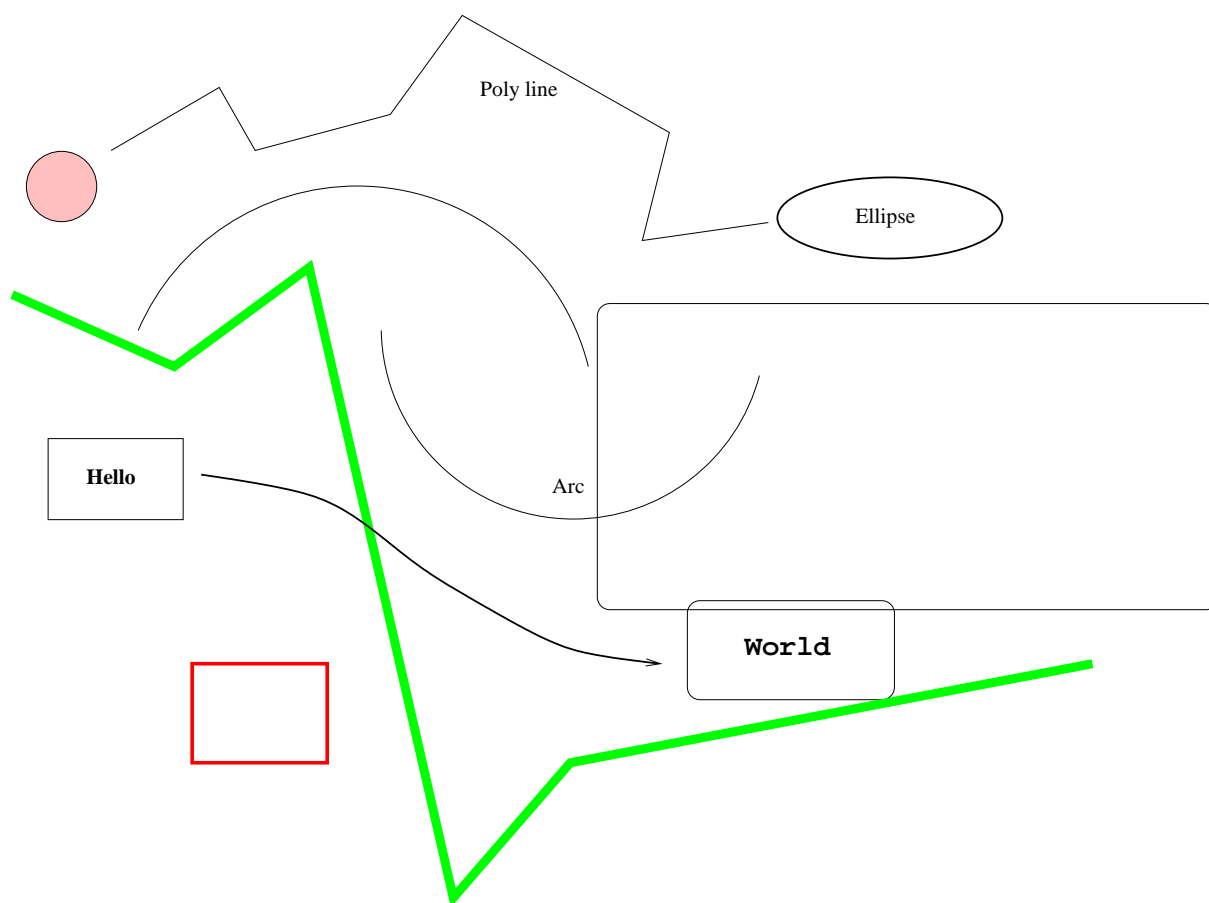


Figure 3.5: Example diagram drawn using `xfig`

Chapter 4

Evaluation

4.1 Printing and binding

Use a “duplex” laser printer that can print on both sides to print two copies of your dissertation. Then bind them, for example using the comb binder in the Computer Laboratory Library.

4.2 Further information

See the Unix Tools notes at

<http://www.cl.cam.ac.uk/teaching/current-1/UnixTools/materials.html>

Chapter 5

Conclusion

I hope that this rough guide to writing a dissertation in L^AT_EX has been helpful and saved you time.

Appendix A

Latex source

A.1 `diss.tex`

A.2 `proposal.tex`

Appendix B

Makefile

B.1 makefile

B.2 refs.bib

Appendix C

Project Proposal

December 30, 2017




Figure C.1: A typical blockchain structure [3]

Project Supervisor: KC Sivaramakrishnan
Director of Studies: Anil Madhavapeddy
Project Overseers: Timothy Jones & Marcelo Fiore

Introduction

The blockchain, in its simplest form, is a tree-like data structure. Chunks of data are stored in 'blocks' which contain the hash of the contents of the previous block. This creates a 'blockchain' which can exhibit branching in the same way that a tree data structure can (see Figure C.1). One of the most important features of a blockchain, is that a change in a block, will alter the block's hash, thereby altering all the future blocks in the chain. This makes it very easy to validate that the data in a blockchain is trustworthy, by verifying the hash in a block, is the same as the hash of it's parent's content.

Blockchain technology has generated a lot of interest in recent times, but mostly in the field of cryptocurrencies. With a simple Proof of Work consensus algorithm, the blockchain can be used to build a secure, distributed ledger of transactions. However, whilst the uses of the blockchain are far wider reaching than cryptocurrencies, progress outside of this field has been much slower.

I will build a pure OCaml, reusable blockchain library to allow the creation of distributed, secure ledgers, which are agreed upon by consensus. The library will allow users to create and add entries to a distributed blockchain ledger with just a few lines of code. The users will also be able to trust that entries in the blockchain are exactly replicated across all nodes in the network.

It will be built on top of Irmin [1] - a distributed database with git-like version control features. Being pure OCaml, the blockchain nodes can be compiled to unikernels or JavaScript to run in the browser. I will evaluate the blockchain by prototyping a decentralised lending library and evaluating the platforms speed and resilience.

Starting point

The project will build upon functionality provided by Irmin [1] which is a distributed database system. Irmin is fast, durable and has the branching capabilities which are required to build a blockchain.

Resources required

I will be using a Macbook provided by OCaml Labs [2] in order to develop the source code for the project. If the Macbook fails, then I will easily be able to transfer my work onto the MCS machines, as my project has no special requirements.

My work will also be backed up to a git repository hosted on GitHub and saved to a dedicated memory stick on a daily basis.

During the evaluation stage I will be running my platform on different cloud based devices and/or Raspberry Pi's. There are many possible providers for cloud computing, including Amazon Web Services and Microsoft Azure. OCaml Labs [2] will provide the necessary funds to acquire these resources.

Background

Consensus

Consensus is a group process where a network of nodes will reach a general agreement. There are different ways of achieving consensus but here are some of the most common:

1. **Proof of Work:** Trust is given to nodes which can prove that they have put in computational work. This is the consensus mechanism used by Bitcoin.
2. **Proof of Stake:** Nodes are selected to validate blocks based on their stake in the blockchain. There are few variations on this algorithm which introduce notions such as delegation or anonymity.
3. **Raft Consensus:** A leader is elected and acts as a governing authority until it fails or disconnects, whereupon a new leader is elected.

Work to be completed

The work for this project will be split into the following major parts.

1. Design and build a module to allow nodes to create and maintain a blockchain ledger. This will include allowing nodes to add blocks to the chain and to form new branches.
2. Design and build a module to allow nodes to interact over a network and to achieve consensus. As highlighted the Background section, there are many different ways to

achieve consensus, and a large part of this work will be to determine which method is most suitable. This decision will take into account a method's failure tolerance in terms of nodes failing and network failure, as well as general speed and any requirements (e.g. computational work for a Proof of Work algorithm).

3. Design an application using these modules. This will take the form of a book lending platform where nodes will be able to register books and lend them to other nodes in the network. This application has been chosen, because the blockchain library should allow for typically centralised applications to be created in a decentralised way. It will also allow for testing of critical features, for example, books should never be 'doubly-spent', i.e. if one user believes they have ownership of a book, then no other user will think the same.
4. Design an evaluation program to simulate different load on the lending platform. This will be run in different configurations in order to measure the performance of the platform.

Evaluation metrics and success criteria

I will consider the project to be a success if the following criteria are achieved:

1. Nodes in the network are able to connect and communicate information.
2. Nodes are able to achieve consensus about the state of the distributed ledger.
3. Nodes are able to reconnect after being individually disconnected.
4. Nodes are able to re-converge after a network partition.

In order to evaluate the performance of the system, I will measure the *throughput* and *speed* of transactions of the book lending platform. Throughput will be measured in transactions per second, and speed will be quantified as time taken to complete a transaction. I will evaluate how these properties vary with respect to the following metrics:

1. **Number of nodes:** I will scale the number of nodes in the network between the range of 2 and 5.
2. **Rate of transactions:** I will vary the number of transactions made per second.

Should I achieve and be able to measure the above criteria within the time frame of my project, I will further test system against the following metrics:

1. **Network latency between nodes**
2. **Network bandwidth of nodes**

Timetable

1. **Michaelmas Weeks 2-4** (12/10/17 - 01/11/17):
Set up an environment for developing OCaml and familiarise myself with the language and it's module system. This is important because the blockchain library needs to be reusable, and therefore well isolated.
2. **Michaelmas Weeks 5-6** (02/11/17 - 15/11/17):
Familiarise myself with Irmin and it's data structures. This is important as I have never used the library before, but it will be used to build the blocks in the blockchain library. In this time I will also begin to design the API of my library.
3. **Michaelmas Weeks 7-8** (16/11/17 - 29/11/17):
Finalise the API and start to build the module for creating and interacting with a distributed ledger. This will also involve investigating which hashing algorithms can be used to form the blockchain data structure.
4. **Christmas Vacation** (30/11/17 - 17/01/18):
Finalise the API of the module for achieving consensus between multiple nodes. This work will also include investigating different methods of consensus and their suitability for my project.
5. **Lent Weeks 1-2** (18/01/17 - 31/01/18):
Build the module for achieving consensus between modules. I will also start work on an lending library application which will be used to evaluate the performance of the blockchain library.
6. **Lent Weeks 3-4** (01/02/18 - 14/02/18):
Finish work on the lending library application and install it on a number of Raspberry Pi and/or cloud based devices. I will also begin work on my dissertation and I aim to complete the Introduction and Preparation chapters.
7. **Lent Weeks 5-6** (05/02/18 - 28/02/18):
Evaluate the performance of the platform by simulating load from each of the devices and measuring the speed of transactions. A stretch goal for this period is also to evaluate a range of further metrics. Additionally I will continue work on my dissertation and aim to complete the Implementation chapter.
8. **Lent Weeks 7-8** (01/03/18 - 14/03/18):
Finish a first draft of my the dissertation by writing the Evaluation and Conclusion chapters. I will also send the dissertation to reviewers to get feedback.
9. **Easter Vacation** (15/03/18 - 25/04/18):
With a first draft of the dissertation completed, I will use this time to review the draft and to make improvements. I will also incorporate feedback from reviewers, and complete the Bibliography and Appendices chapters.

10. **Easter Weeks 1-2** (26/04/18 - 09/05/18):

Conclude work on dissertation by incorporating final feedback from reviewers.

11. **Easter Week 3-Submission Deadline** (10/05/18 - 08/05/18):

I aim to have completed the dissertation by this point, and to be focusing on my studies. However, this time may be needed to make any final changes.

Bibliography

- [1] Irmin - A pure OCaml, distributed database that follows the same design principles as Git.
<https://github.com/mirage/irmin>
- [2] OCaml Labs - An initiative based in the Computer Laboratory to promote research, growth and collaboration within the wider OCaml community
<http://ocamlabs.io/>
- [3] Image of blockchain data structure from Wiki Commons.
<https://commons.wikimedia.org/wiki/File:Blockchain.png>